

Numerical Optimization and its Applications: Final Project

Shahaf Finder
204037618

Roy Uziel
203398854

July 10, 2018


1 Objective Loss Function

- We've computed Softmax and Loss in softmax.m and loss.m.
- The derivatives for the Softmax are in loss_grad_theta.m (gradient by weights) and loss_grad_x.m (gradient by X).
- test_loss.m is the gradient test for the loss function.
- loss_SGD.m is an SGD variant for the loss function.
- Q3.m is a script used to run the SGD.


We didn't notice a significant difference using different batch size and learning rate.

Results for learning rate: 0.005, batch size: 100

for validation data:



for training data:



2 Residual Neural Network

2.1 Jacobian Transpose Times Vector

`ResNN_jac_theta_t_mul.m` and `ResNN_jac_x_t_mul.m` are the jacobian transpose times vector for each layer.

Both can receive multiple samples X (as an n times m matrix), and their respective V (as an n times m matrix) to multiply by.

- `ResNN_jac_theta_t_mul.m` returns the average result among all samples.
- `ResNN_jac_x_t_mul.m` returns the multiplication of each sample with it's v (as an n times m matrix).

In order to test those functions we've created two auxiliary functions `ResNN_jac_theta_mul.m` and `ResNN_jac_x_mul.m`.

We used the gradient test to make sure these functions are correct, and then used the jacobian transpose test with the required functions.

Those tests are in `test_resnn.m`.

2.2 Forward Pass and Back Propagation

`forward_pass.m` and `back_propagation.m` are the functions required.

To test them we created `test_net.m` that uses forward pass and back propagation to compute the gradient, then each iteration we used forward pass to compute the loss value after moving to a direction.

2.3 SGD and Conclusions

ResNN_SGD.m is the SGD variant for the entire network.

run_net.m is a script used to run the SGD.

We've tested different network lengths and batch sizes on different data sets. We got to the following conclusions:

- Increasing batch size increased the training time, and above 50 it had little effect on results.
- Each dataset required different network lengths for good results.
- Lowering the learning rate resulted in a more steady but slower convergence. We found out that, generally, 0.01 is a good balance between speed and steadiness.

3 Improving on SGD

In seek to improve the results of SGD, we tried the following:

- Decreasing alpha according to iteration
After 100 iterations we calculated a new alpha by: $learning_rate * 5 / \sqrt{i}$.
This is in order to half the learning rate and keep decreasing as the iterations continue.
The method gave minor improvement.
- Adding momentum (as seen in class)
The method got much better results in convergence speed, but was unsteady in the long run.
So in order to deal with the unsteadiness, after 100 iterations we calculated a new gamma using: $moment_gamma * 10 / \sqrt{i}$.
This decreases the momentum as iterations continue.
The results were a fast convergence in a more steady way.
- Taking only the rejection part of the momentum.
While trying to optimize even further, we tried to take only the projection of the momentum over the gradient.
We calculated it by: $m_proj = g * m' * g / (norm(g)^2)$.
We found out that by adding only the projection part, the results got worse, so we tried adding the rejection part.
Calculated by: $m_rej = m - m_proj$.
The results were usually better, or at worse - the same, than the normal momentum.

Using these 3 improvements made us able to reduce the number of layers without hurting the results too much.

Comparison between the 3 improvements.


Blue - using only decreasing alpha.

Orange - using decreasing alpha, and decreasing momentum.


Yellow - using all of the improvements specified above.

GMM

batch size: 50, learning rate: 0.01, number of layers: 1, gamma(only for momentum): 0.7




batch size: 50, learning rate: 0.01, number of layers: 2, gamma(only for momentum): 0.7




Peaks

batch size: 50, learning rate: 0.01, number of layers: 4, gamma(only for momentum): 0.7




batch size: 50, learning rate: 0.01, number of layers: 6, gamma(only for momentum): 0.7



SwissRoll

batch size: 50, learning rate: 0.01, number of layers: 4, gamma(only for momentum): 0.7



batch size: 50, learning rate: 0.01, number of layers: 5, gamma(only for momentum): 0.7

